

# State-of-the-Art Evaluation and Benchmarking for LLM and Multi-Agent Systems

## The Evolving Landscape of AI System Evaluation

The rapid transition of Large Language Models (LLMs) from research prototypes to core components of enterprise applications has created an urgent and complex challenge: how to rigorously evaluate their performance, reliability, and safety. As these models are increasingly embedded into autonomous or semi-autonomous agents that reason, plan, and interact with external systems, traditional evaluation paradigms have proven insufficient. A fundamental shift in methodology is required, moving from static assessments of text generation to dynamic, holistic evaluations of complex, interactive systems. This new landscape demands a synthesis of principles from natural language processing (NLP), human-computer interaction (HCI), and software engineering to address the unique challenges posed by probabilistic, emergent agentic behavior, especially within the stringent context of enterprise deployment.

## From Static Models to Dynamic Agents: A Paradigm Shift in Evaluation

The evaluation of a standalone LLM is fundamentally different from the evaluation of an LLM-powered agent. An LLM is typically assessed on its capacity for text generation, summarization, or question-answering in a controlled, static environment. In this context, evaluation is akin to testing a component in isolation. A compelling analogy frames this distinction clearly: "LLM evaluation is like examining the performance of an engine. In contrast, agent evaluation assesses a car's performance comprehensively, as well as under various driving conditions".<sup>1</sup>

LLM agents operate in dynamic, interactive environments where they must perform a sequence of actions to achieve a goal. This process involves sophisticated capabilities that

are not captured by traditional NLP metrics. Agents reason about tasks, formulate multi-step plans, execute tools, leverage memory to maintain context over long horizons, and even collaborate with other agents or humans.<sup>1</sup> The success of an agent is not merely determined by the quality of a single text output, but by the cumulative outcome of its actions and its ability to navigate a complex, often unpredictable environment.

This shift introduces a new layer of complexity rooted in the inherent nature of agentic systems. Unlike traditional software, which exhibits deterministic and static behavior, LLM agents are probabilistic. The same input may not always produce the exact same sequence of actions, making traditional software testing methods inadequate.<sup>1</sup> Evaluating these systems requires a new framework that can account for this dynamic behavior and assess the entire lifecycle of a task, from initial planning to final execution. This evolution in system capability mirrors the historical progression of software testing methodologies. Early software development relied on unit tests to validate individual functions in isolation, much like how early LLM benchmarks tested a model's raw text generation. As software systems grew into complex, interconnected microservices, the focus shifted to integration and end-to-end (E2E) testing to validate the entire workflow. Similarly, the components of an agent—its planning module, memory, and tool-use functions—act as interconnected services. A comprehensive evaluation must therefore encompass not only "unit tests" for the core model's capabilities but also "integration tests" for how these components interact and "E2E tests" that measure the agent's success in achieving a user's ultimate goal.

## A Taxonomy of Modern Evaluation: Objectives and Processes

To bring clarity to the fragmented landscape of agent evaluation, a structured, two-dimensional taxonomy provides a robust framework for systematic assessment. This framework, proposed by Mohammadi et al. (2025), organizes evaluation methodologies along two primary axes: **Evaluation Objectives** (what to evaluate) and the **Evaluation Process** (how to evaluate).<sup>1</sup>

### Dimension 1: Evaluation Objectives

This dimension focuses on the specific targets of evaluation, breaking down an agent's performance into distinct, measurable categories.

- **Agent Behavior:** This category is outcome-oriented, focusing on the end results of an agent's actions. It includes metrics such as task completion success rate and the overall

quality of the final output. Output quality is a multifaceted concept encompassing accuracy, relevance, clarity, coherence, and adherence to task requirements or predefined constraints.<sup>1</sup> This objective answers the fundamental question: "Did the agent successfully achieve the user's goal?"

- **Agent Capabilities:** This category is process-oriented, examining the underlying competencies that enable an agent to function. It provides insight into how an agent achieves its goals. Key capabilities include:
  - **Planning and Reasoning:** The ability to decompose a complex goal into a logical sequence of actionable steps.<sup>1</sup>
  - **Tool Use:** The capacity to correctly select and invoke external tools or functions, including generating the appropriate parameters for execution.<sup>1</sup>
  - **Memory and Context Retention:** The ability to maintain and recall relevant information across multi-turn interactions to inform future decisions.<sup>1</sup>
  - **Multi-Agent Collaboration:** The effectiveness of communication, coordination, and synchronization when working with other agents to achieve a shared objective.<sup>1</sup>
- **Reliability:** This objective assesses the consistency and robustness of an agent's behavior. It measures whether an agent produces consistent outcomes for the same input and how robustly it performs when inputs vary or when it encounters unexpected errors or environmental changes.<sup>1</sup>
- **Safety and Alignment:** This critical objective evaluates an agent's trustworthiness and security. It encompasses fairness (absence of bias), compliance with regulations and policies, and the prevention of harmful, unethical, or malicious behaviors.<sup>1</sup>

## Dimension 2: Evaluation Process

This dimension describes the methodologies and resources used to conduct the evaluation.

- **Interaction Mode:** This distinguishes between two primary modes of assessment. **Static evaluation** involves providing the agent with a fixed set of inputs and assessing its responses, typical of traditional benchmarks. **Interactive assessment**, in contrast, involves dynamic engagement where the agent interacts with a user, another agent, or a simulated environment, allowing for the evaluation of its performance in more realistic, multi-turn scenarios.<sup>1</sup>
- **Evaluation Data:** This component concerns the datasets and benchmarks used for evaluation. It covers both **synthetic data**, which is artificially generated to test specific capabilities, and **real-world data**, which is collected from actual user interactions or applications. This category also includes the growing number of specialized benchmarks tailored to specific domains like software engineering, finance, or healthcare.<sup>1</sup> The increasing complexity of agentic systems also necessitates that evaluation processes

account for multilingual and multimodal inputs, requiring tailored metrics and strategies to assess performance across these variations.<sup>1</sup>

## Key Challenges in Enterprise Agent Evaluation

While academic benchmarks provide a crucial foundation for measuring agent performance, deploying these systems in an enterprise context introduces a set of challenges that are often overlooked by standard evaluation frameworks.<sup>1</sup> These challenges highlight a significant gap between performance in a controlled lab environment and readiness for production, a gap defined less by raw capability and more by governance, trust, and security. An agent can achieve a state-of-the-art score on a public leaderboard yet be entirely undeployable in a regulated industry like finance or healthcare due to its failure to meet these enterprise-grade requirements.

This reality necessitates that a comprehensive evaluation strategy must extend beyond public benchmarks to include an internal layer of testing focused on these critical, non-functional requirements. This transforms agent evaluation from a purely technical MLOps task into a cross-functional responsibility involving legal, compliance, and security stakeholders.

The primary enterprise-specific challenges include:

- **Security and Role-Based Access Control:** Enterprise agents often need to interact with sensitive internal data and proprietary systems. A critical evaluation requirement is to ensure the agent strictly adheres to role-based access control (RBAC) policies, preventing unauthorized data access or actions.<sup>2</sup>
- **Reliability Guarantees and Auditing:** In many industries, actions taken by an agent must be auditable and highly reliable. Evaluation must provide guarantees that the agent behaves predictably and consistently, especially for mission-critical tasks. This includes robust logging and traceability to support compliance and audit trails.<sup>1</sup>
- **Long-Horizon and Dynamic Interactions:** Enterprise workflows are often complex and long-running, requiring agents to maintain context and adapt over extended, multi-turn interactions. Evaluating an agent's performance in these long-horizon scenarios is significantly more challenging than assessing single-turn tasks.<sup>2</sup>
- **Compliance and Regulatory Adherence:** Agents operating in regulated domains must comply with a host of legal and industry-specific standards. Evaluation frameworks must be designed to verify that an agent's behavior and outputs consistently adhere to these rules, a requirement rarely addressed in academic benchmarks.<sup>2</sup>

# A Comparative Analysis of State-of-the-Art Benchmarking Suites

The limitations of older benchmarks, which often suffer from data contamination and a narrow focus on simple tasks, have spurred the development of a new generation of specialized evaluation suites. These modern benchmarks are designed to test the advanced capabilities of today's LLMs and multi-agent systems with greater rigor and realism. This trend marks a significant maturation of the field, moving away from the pursuit of a single, monolithic benchmark toward the adoption of a portfolio of specialized tools. Just as a software engineering team uses different testing frameworks for its database, user interface, and backend logic, an AI development team must now select the right benchmark to evaluate a specific capability, be it multilingual code generation, complex planning, or general cognitive reasoning. This portfolio-based approach enables a more nuanced and accurate assessment of an agent's true capabilities.

A second, powerful trend is the use of LLMs themselves to bootstrap the creation of more challenging and higher-quality evaluation datasets. By automating or augmenting the labor-intensive process of creating test cases, frameworks like AutoCodeBench and BenchMAX are accelerating the entire field's rate of progress. This creates a virtuous cycle where more capable models help build better evaluation tools, which in turn drive the development of even more capable models.

## Code Generation and Reasoning: AutoCodeBench vs. LiveCodeBench

The ability to generate and reason about code is a critical measure of an LLM's advanced capabilities. Two recent benchmarks, LiveCodeBench and AutoCodeBench, have emerged as state-of-the-art tools in this domain, each offering a unique and powerful methodology.

### LiveCodeBench

LiveCodeBench is a holistic and contamination-free benchmark designed to provide a realistic measure of a model's coding abilities on unseen problems.<sup>4</sup>

- **Methodology:** Its core innovation is the continuous collection of new, high-quality coding problems from competitive programming platforms like LeetCode and AtCoder.

Crucially, each problem is annotated with its release date. This allows for temporal-based evaluation, where a model with a known training cutoff date can be tested exclusively on problems published after that date, providing a true measure of its generalization ability and effectively mitigating the issue of training data contamination.<sup>4</sup>

- **Evaluation Scenarios:** LiveCodeBench moves beyond simple code generation to assess a broader suite of code-related skills. Its evaluation scenarios include:
  - **Code Generation:** The standard task of producing a functional code solution.
  - **Self-Repair:** The agent's ability to debug and fix its own incorrect code.
  - **Test Output Prediction:** A reasoning task where the model must predict the output of a given code snippet and test case.
  - **Code Execution:** The ability to mentally simulate or execute code to determine its outcome.<sup>4</sup>
- **Key Findings:** By testing on a continuously updated set of problems, LiveCodeBench has revealed potential overfitting issues with older benchmarks. It has shown that some models achieve high scores on static benchmarks like HumanEval but exhibit a significant performance drop on newer, unseen problems, suggesting their performance may be inflated due to data contamination.<sup>4</sup>

## AutoCodeBench

AutoCodeBench introduces a novel, scalable, and fully automated approach to creating a large-scale, multilingual code generation benchmark.<sup>5</sup>

- **Methodology:** At the heart of AutoCodeBench is AutoCodeGen, an automated workflow that leverages an LLM-sandbox interaction. Instead of relying on manual annotation, the process uses an LLM to generate code solutions, test inputs, and, in a reverse-order process, the programming problems themselves. The correctness of the generated code and tests is verified at each step by executing them in a secure, multilingual sandbox that supports over 20 programming languages.<sup>5</sup> The resulting dataset undergoes multiple filtering stages to ensure high quality, difficulty, and diversity.
- **Multilingual Focus:** A key strength of AutoCodeBench is its extensive multilingual coverage. The final benchmark contains 3,920 challenging problems evenly distributed across 20 different programming languages, making it one of the most comprehensive benchmarks for assessing multilingual coding proficiency.<sup>5</sup>
- **Benchmark Versions:** To cater to different evaluation needs, the project offers several versions:
  - **AutoCodeBench:** The full, large-scale benchmark designed to measure absolute multilingual performance.
  - **AutoCodeBench-Lite:** A smaller subset of 1,586 problems, designed to be more efficient for measuring relative performance differences between models.

- **AutoCodeBench-Complete:** A version tailored for evaluating the few-shot capabilities of base models using a code completion format.<sup>5</sup>

## Comparative Analysis

LiveCodeBench and AutoCodeBench represent two distinct but complementary philosophies in modern code evaluation. LiveCodeBench's strength lies in its rigorous approach to preventing data contamination and its holistic assessment of multiple coding-related skills. It is the ideal tool for verifying a model's true generalization ability. AutoCodeBench's innovation lies in its scalability and multilingual breadth, enabled by its fully automated generation pipeline. It provides an unparalleled resource for teams needing to evaluate performance across a wide array of programming languages, particularly those that are less common in other benchmarks.

## Assessing General Cognitive Abilities: The Role of AGIEval

While code generation is a powerful proxy for reasoning, a broader measure of a model's cognitive abilities is often required. AGIEval addresses this need by creating a human-centric benchmark grounded in high-stakes examinations designed to test human intellect.<sup>9</sup>

- **Methodology:** AGIEval is derived from 20 official, public, and high-standard admission and qualification exams from around the world. These include general college admission tests like the American SAT and the Chinese Gaokao, as well as specialized exams for law school admission, mathematics competitions, and civil service qualifications.<sup>9</sup> By using tasks that are standardized for human evaluation, AGIEval provides a stable and well-understood baseline for assessing an LLM's general problem-solving and reasoning skills.
- **Tasks and Data Format:** The benchmark consists of 18 multi-choice question (MCQ) tasks and two cloze (fill-in-the-blank) tasks. The data is provided in a simple, structured JSON format that includes fields for a passage (if applicable), the question, options, and the correct label or answer. This standardized format makes it straightforward to integrate AGIEval into automated evaluation pipelines.<sup>9</sup>
- **Significance:** AGIEval serves as a robust proxy for a model's general intelligence in domains that require a combination of knowledge, logic, and reasoning. It moves evaluation beyond narrow NLP tasks and toward a more holistic assessment of the cognitive abilities that are critical for advanced agentic systems.

## Evaluating Multi-Agent Collaboration: Planning and Scheduling with REALM-Bench

As AI systems evolve from single agents to collaborative multi-agent systems (MAS), new benchmarks are needed to evaluate their ability to coordinate and plan effectively. REALM-Bench is a comprehensive suite designed specifically for this purpose, focusing on real-world, dynamic planning and scheduling tasks.<sup>10</sup>

- **Methodology:** REALM-Bench is an extensible evaluation suite that tests the planning, coordination, and adaptation capabilities of both individual LLMs and multi-agent systems. Its distinguishing feature is a systematic focus on system reliability in the face of unexpected events and dynamic disruptions, such as resource shortages or environmental changes that require real-time replanning.<sup>12</sup>
- **Key Features:** The benchmark is designed to be highly realistic and challenging. It incorporates:
  - **Multi-agent Coordination:** Tasks that require multiple agents to communicate and synchronize their actions.
  - **Inter-agent Dependencies:** Scenarios where the actions of one agent create constraints or opportunities for others.
  - **Dynamic Disruptions:** Unexpected events that force the system to adapt and replan in real-time.
  - **Scalable Complexity:** Each problem can be scaled along three dimensions: the number of parallel planning threads (agents), the complexity of inter-dependencies, and the frequency of disruptions.<sup>10</sup>
- **Problem Series:** The benchmark is organized into two complementary categories:
  - **P-series:** Ten problems covering general real-world planning scenarios, such as urban ride-sharing, disaster relief logistics, and global supply chain management.
  - **J-series:** Four problems based on the canonical Job-Shop Scheduling Problem (JSSP), a well-studied class of combinatorial optimization problems. These problems are arranged in five tiers of progressing difficulty, starting from simple, single-agent static planning and culminating in large-scale, dynamic multi-domain integration.<sup>12</sup> REALM-Bench is an essential tool for any team building multi-agent systems, as it provides a standardized way to measure the effectiveness of their orchestration and collaboration logic.

## Gauging Global Proficiency: Advanced Multilingual Assessment with



# BenchMAX

For AI agents intended for a global audience, the ability to perform complex tasks consistently across multiple languages is paramount. BenchMAX is a state-of-the-art benchmark designed to evaluate these advanced, language-agnostic capabilities, moving far beyond simple translation or text understanding.<sup>13</sup>

- **Methodology:** BenchMAX is a comprehensive, multi-way multilingual benchmark that assesses six crucial capabilities across 17 diverse languages, including English, Spanish, German, Japanese, Chinese, Swahili, and Arabic.<sup>13</sup>
- **Evaluated Capabilities:** Unlike previous multilingual benchmarks that focused on simple understanding tasks, BenchMAX evaluates the advanced skills required by modern agents:
  - Instruction Following
  - Reasoning (Math and Science)
  - Code Generation
  - Long Context Understanding
  - Tool Use
  - Translation (General and Domain-Specific) <sup>13</sup>
- **Quality Assurance:** To ensure the high quality and cultural appropriateness of its datasets, BenchMAX employs a rigorous, multi-stage construction pipeline. The process begins with machine translation of English source data. Each translated sample is then independently reviewed and post-edited by three distinct native-speaking annotators. Finally, a superior LLM is used as a judge to select the best and most accurate translation, ensuring the final dataset is of extremely high quality.<sup>13</sup>
- **Significance:** BenchMAX provides a critical tool for understanding the true multilingual proficiency of LLMs. Its extensive experiments have revealed that simply scaling up a model's size does not necessarily close the performance gap between English and other languages, highlighting the need for more targeted multilingual training and evaluation.<sup>13</sup> For any organization deploying a global AI product, BenchMAX offers an essential framework for ensuring equitable and high-quality performance for all users.

The following table provides a consolidated comparison of the key benchmarks discussed, offering a strategic overview to guide selection based on specific evaluation needs.

Bench mark Name	Primary Focus	Key Method ology	Evaluat ed Capabili ties	Multilin gual Support	Contam ination Strateg y	Strengt hs	Weaknes ses/Li mitation s
LiveCo	Contam	Continu	Code	Primaril	Tempor	Excelle	Smaller

<b>deBench</b>	ination-free code evaluation	ous collection of new problems from coding sites with release-date annotation.	Generation, Self-Repair, Test Output Prediction, Code Execution.	y English, but problems can be in any language.	al cutoff based on problem release dates.	nt for measuring true generalization; holistic skill assessment.	dataset size compared to automated methods; less language diversity.
<b>AutoCodeBench</b>	Scalable, multilingual code generation	Fully automated AutoCodeGen workflow using LLM-Sandbox interaction and reverse problem generation.	Code Generation (across 20 languages).	<b>High (20 languages)</b> , evenly distributed.	Data filtering and diversity sampling.	Massive scale and language breadth; fully automated and human-free.	Does not evaluate broader skills like self-repair; potential for LLM-induced biases.
<b>AGIEval</b>	General cognitive and reasoning abilities	Derived from 20 high-standard human exams (SAT, Gaokao, etc.).	Human-level problem-solving, logic, math, and reading	<b>High (English &amp; Chinese)</b> tasks included.	Uses established, static human exams.	Provides a stable, human-grounded measure of general	Static dataset; does not test interactive or agentic capabilities.

			compre hension .			intellige nce.	
<b>REALM -Bench</b>	Multi-a gent plannin g and coordin ation	Extensi ble suite of dynami c plannin g and schedul ing tasks with disrupti ons.	Multi-a gent plannin g, coordin ation, adaptat ion, reactive replanni ng.	Langua ge-agn ostic (proble m descrip tions).	Not applica ble (focus is on dynami c interact ion).	The premier benchm ark for multi-a gent systems ; tests reliabilit y under stress.	Highly speciali zed; not intende d for general LLM capabili ty testing.
<b>Bench MAX</b>	Advanc ed multilin gual capabili ties	Multi-w ay parallel benchm ark with a rigorou s human- in-the-l oop translati on and annotat ion pipeline .	Instruct ion Followin g, Reasoni ng, Code Generat ion, Tool Use, Long Context , Translat ion.	<b>Very High (17 langua ges).</b>	High-q uality human annotat ion and review.	The most compre hensive benchm ark for advanc ed multilin gual skills.	Dataset constru ction is resourc e-inten sive; relies on LLM-as -judge for final selectio n.
<b>SWE-B ench</b>	Real-w orld softwar e enginee	Source d from real GitHub issues	Agentic proble m-solvi ng, reposit	Primaril y English -based codeba	Uses real, historic al data.	Highly realistic measur e of practic	Narrowl y focused on issue

	ring tasks	in popular reposit ories.	ory-leve l code underst anding, resolvin g real-wo rld bugs.	ses.		al softwar e enginee ring ability.	resoluti on; comple x setup require d.
--	---------------	------------------------------------	--	------	--	---	---

## Integrating Continuous Evaluation into Modern Development Workflows (CI/CD)

The maturation of AI development from an experimental science into a robust engineering discipline hinges on the adoption of principles from modern software development, most notably Continuous Integration and Continuous Deployment (CI/CD). By integrating automated evaluation directly into the development lifecycle, teams can treat changes to AI systems—whether to the underlying model, a prompt template, or an agent's toolset—with the same rigor as changes to traditional application code. This practice transforms evaluation from a sporadic, manual process into a continuous, automated safety net that catches regressions, validates improvements, and ensures a consistent level of quality and reliability for production systems. This shift has profound organizational implications, requiring close collaboration between ML scientists, software engineers, and MLOps specialists, and necessitates the versioning of all AI system components, including prompts and evaluation datasets, as part of an "infrastructure-as-code" philosophy.

### Principles of Continuous Evaluation for LLM-Powered Systems

Continuous evaluation is the practice of automatically running a suite of tests and benchmarks whenever a change is made to an AI system. This creates a rapid feedback loop that allows developers to identify and fix issues early, before they reach production.<sup>17</sup>

A successful continuous evaluation strategy requires careful planning around several key decisions:

- **What to Measure:** The evaluation suite should include a portfolio of metrics that cover the system's key performance indicators (KPIs), such as task success rate, output quality (e.g., factual consistency, relevance), and operational metrics (e.g., latency, cost).<sup>19</sup>
- **When to Measure:** The frequency of evaluation depends on the specific trigger. A lightweight suite of critical tests might run on every code commit or pull request, while a more comprehensive and resource-intensive benchmark suite might be scheduled to run nightly or weekly.<sup>19</sup>
- **How to Manage Resources:** LLM evaluations can be computationally expensive and time-consuming. The strategy must balance the need for thorough testing with constraints on budget and development velocity. This may involve using smaller, targeted test sets for frequent checks and reserving full benchmark runs for less frequent, scheduled pipelines.<sup>19</sup>

## Pattern 1: GitHub Actions for Automated Benchmark Execution

GitHub Actions provides a powerful and flexible platform for automating CI/CD workflows, including the continuous evaluation of LLM agents. Workflows are defined in YAML files located in the `.github/workflows/` directory of a repository and can be triggered by events such as a `pull_request` or a `schedule`.<sup>21</sup>

### Reproducible Example: Continuous Evaluation with AutoCodeBench

The following example demonstrates a GitHub Actions workflow that automatically evaluates a code generation model using the AutoCodeBench framework. This workflow is triggered whenever a pull request is opened against the main branch. It sets up the necessary environment, including the multilingual sandbox, runs the evaluation, and uses the `github-action-benchmark` tool to compare the new result against the baseline from the main branch, failing the check if a significant performance regression is detected.

**File:** `.github/workflows/autocodebench-eval.yml`

YAML

name: 'AutoCodeBench Evaluation'

on:

pull\_request:

branches: [ main ]

permissions:

contents: read

pull-requests: write

checks: write

jobs:

evaluate-model:

runs-on: ubuntu-latest

services:

sandbox:

image: hunyuansandbox/multi-language-sandbox:v1

ports:

- 8080:8080

options: --cap-add=NET\_ADMIN

steps:

- name: Checkout repository

uses: actions/checkout@v4

- name: Set up Python

uses: actions/setup-python@v5

with:

python-version: '3.10'

- name: Install dependencies

run: |

python -m pip install --upgrade pip

pip install -r requirements.txt # Assuming requirements.txt contains necessary libraries

- name: Generate model outputs

# This step simulates generating code solutions from the model being tested.

# In a real pipeline, this would involve calling the model API.

# For this example, we assume the outputs are pre-generated in 'model\_output.jsonl'.

run: |

echo "Simulating model inference..."

# python run\_inference.py --model-path \${ secrets.MODEL\_ENDPOINT } --output-file  
model\_output.jsonl

```

- name: Run AutoCodeBench Evaluation
  id: run_eval
  # This script calls the sandbox service to evaluate the model's outputs.
  # It must be configured to output the final score in a format github-action-benchmark can parse.
  run: |
    python call_sandbox.py \
      --input_file model_output.jsonl \
      --output evaluation_results.jsonl \
      --server_ip localhost \
      --server_port 8080 \
      --concurrency 32

    # Extract the pass@1 score and format it for the benchmark action
    PASS_RATE=$(python -c "import json; data=json.load(open('evaluation_results.jsonl'));
print(data['summary']['pass@1'])")
    echo " > benchmark_output.json

- name: Store benchmark results
  uses: rhyds/github-action-benchmark@v1
  with:
    name: AutoCodeBench Performance
    tool: 'customSmallerIsBetter' # Placeholder, actual comparison logic is below
    output-file-path: benchmark_output.json
    github-token: ${{ secrets.GITHUB_TOKEN }}
    # Compare results and add a comment to the PR
    comment-on-alert: true
    # Fail the workflow if the new result is 10% worse than the previous one
    alert-threshold: '110%'
    fail-on-alert: true

```

This workflow demonstrates a robust pattern for continuous evaluation <sup>8</sup>:

1. **Environment Setup:** It uses a service container to run the required multilingual sandbox, ensuring the evaluation environment is consistent and isolated.<sup>21</sup>
2. **Execution:** It runs the evaluation script provided by the AutoCodeBench repository, which communicates with the sandbox service to score the model's outputs.<sup>8</sup>
3. **Regression Testing:** It leverages a third-party action to store, compare, and alert on benchmark results. This automatically prevents merging changes that cause a significant drop in performance, enforcing a quality gate for the AI system.<sup>23</sup>

## Pattern 2: GitLab CI for Scheduled and Trigger-Based Agent Validation

GitLab CI/CD provides another powerful ecosystem for automating evaluation pipelines. Configuration is managed through a `.gitlab-ci.yml` file at the root of the repository, which defines jobs, stages, and rules for execution.<sup>26</sup>

## Reproducible Example: Nightly Evaluation with AGIEval

The following example illustrates a GitLab CI pipeline that performs a nightly evaluation of an LLM using the AGIEval benchmark. This is useful for monitoring model performance over time or for running more extensive tests that are too slow for a per-commit workflow.

**File: `.gitlab-ci.yml`**

YAML

stages:

- evaluate
- report

variables:

```
# Store the OpenAI API key as a masked CI/CD variable in the project settings
# OPENAI_API_KEY: "your_api_key_here"
PYTHON_VERSION: "3.10"
```

nightly-agieval-evaluation:

```
stage: evaluate
image: python:${PYTHON_VERSION}
```

before\_script:

- pip install -r requirements.txt # Assuming requirements for AGIEval scripts

script:

- echo "Starting AGIEval evaluation for model..."
- ```
# Set the API key for the script to use
```
- export OPENAI\_API\_KEY=\$OPENAI\_API\_KEY



```
# Run the prediction script from the AGIEval repository
```

```
- python run_prediction.py --model gpt-4o --output_dir./results
```

```
# Run the post-processing and evaluation script
```

```
- python post_process_and_evaluation.py --input_dir./results --output_file agieval_scores.json
```

```
artifacts:
```

```
paths:
```

```
- agieval_scores.json
```

```
when: on_success
```

```
expire_in: 7 days
```

```
rules:
```

```
# Run this job only on a nightly schedule
```

```
- if: '$CI_PIPELINE_SOURCE == "schedule"'
```

This pipeline demonstrates key GitLab CI features for continuous evaluation <sup>27</sup>:

1. **Stages:** The pipeline is organized into logical stages (evaluate, report), ensuring a clear and ordered execution flow.<sup>27</sup>
2. **Docker Image:** It specifies a Docker image (python:3.10) to ensure a consistent and reproducible execution environment for the evaluation scripts.<sup>27</sup>
3. **Secrets Management:** It securely handles the OPENAI\_API\_KEY by leveraging GitLab's built-in CI/CD variables, which should be configured as "masked" in the project's settings to prevent exposure in job logs.<sup>30</sup>
4. **Artifacts:** The final evaluation report (agieval\_scores.json) is saved as a job artifact, making it easily accessible for review and downstream processing. The artifact is set to expire after 7 days to manage storage.<sup>26</sup>
5. **Scheduling:** The rules keyword ensures that this job only runs when triggered by a pipeline schedule (e.g., configured to run every night at midnight), separating it from the regular development workflow.<sup>26</sup>

## Advanced Methodologies for Agentic Reasoning and Workflow Verification

While end-to-end benchmarks are essential for measuring overall task success, they often treat the agent as a "black box," revealing *if* it failed but not *why*. To build truly reliable and trustworthy agentic systems, it is critical to move beyond outcome-based evaluation and implement methodologies that can verify the internal processes of the agent. This involves

inspecting the agent's reasoning chain, validating its tool usage, and diagnosing the precise points of failure within its workflow. This shift toward process-level verification is arguably the most significant step in the maturation of agent evaluation. It is analogous to the role of code reviews and static analysis in traditional software engineering; the focus is not just on the final output of the program but on the quality and correctness of the underlying logic. For agentic systems, the "source code" is the reasoning trace. Consequently, robust logging and observability are not optional features but absolute prerequisites for any advanced evaluation, making tracing platforms a non-negotiable component of the modern MLOps stack for AI agents.

## Implementing Plan→Apply Evaluation Pipelines

A structured and scalable approach to agent evaluation can be modeled as a "Plan→Apply" or "Fetch→Run→Save" pipeline. This pattern provides a systematic way to conduct offline, incremental evaluations on data captured from a production or staging environment, creating a continuous feedback loop for improvement.<sup>31</sup>

1. **Plan (Fetch & Define):** The first step is to define the scope and criteria for the evaluation. This involves programmatically fetching traces of agent interactions from an observability platform or logging database. These traces can be filtered by date, user tags, or specific events to create a targeted dataset for analysis. Once the data is fetched, the specific evaluation metrics are defined. These can range from simple, code-based checks to complex, qualitative assessments.<sup>31</sup>
2. **Apply (Run Evaluations):** The core of the pipeline is the execution of the evaluation logic on the fetched traces. This step can employ a variety of "scorers":
  - **Deterministic Scorers:** These are code-based checks for objective criteria. For example, a scorer can verify if an agent's output is valid JSON, if a SQL query it generated is syntactically correct, or if a final answer matches a known ground truth value.<sup>32</sup>
  - **LLM-as-a-Judge:** For more nuanced, qualitative metrics like relevance, coherence, or helpfulness, another powerful LLM can be used as an evaluator. Frameworks like DeepEval's G-Eval or custom prompting strategies can be used to score an agent's output against a predefined rubric.<sup>34</sup> For example, an evaluation prompt might ask an LLM judge to rate the "joyfulness" of a response on a numerical scale or categorize its "tone".<sup>31</sup>
3. **Verify (Save Results):** The final step is to persist the evaluation scores, attaching them back to the original traces in the observability platform. This creates a rich, auditable record of the agent's performance over time, allowing teams to track regressions, measure the impact of changes, and identify areas for improvement.<sup>31</sup> This entire pipeline can be orchestrated and scheduled (e.g., using a cron job) to run periodically, ensuring

continuous monitoring of the agent's quality in production.<sup>31</sup>

## Verifying Multi-Step Reasoning and Tool Use Integrity

For agents that perform complex, multi-step reasoning, such as those employing a Chain-of-Thought (CoT) approach, verifying the final answer is insufficient. A correct answer derived from flawed logic is a "lucky guess" and cannot be trusted in high-stakes applications. Therefore, it is essential to verify the integrity of each step in the reasoning process.<sup>37</sup>

A useful taxonomy for analyzing multi-step reasoning involves three phases: **Generate**, **Evaluate**, and **Control**.<sup>38</sup> The evaluation phase is critical for verification and can be implemented using several techniques:

- **Self-Assessment:** The agent itself is prompted to review and critique its own reasoning steps. This can involve checking for logical consistency, identifying factual errors, or confirming that each step contributes to solving the problem.<sup>37</sup> The Chain of Verification (CoVe) method is a formalization of this approach, where the model explicitly plans, executes, and then verifies its own work.<sup>37</sup>
- **Tool-Based Evaluation:** When a reasoning step involves a deterministic operation (e.g., a mathematical calculation or a code execution), an external tool can be used to verify its correctness. For example, if an agent reasons that "25\*4=100", a simple Python interpreter can be called to validate this step independently.<sup>38</sup>
- **Deductive Verification:** This technique involves decomposing a complex reasoning chain into its fundamental, step-by-step logical deductions. Each individual deduction can then be verified for validity, ensuring that the entire chain is sound from start to finish. This approach is particularly powerful for identifying subtle logical fallacies or hallucinations that might otherwise go unnoticed.<sup>39</sup>

To further enhance the structure and verifiability of complex reasoning, advanced frameworks like **Agentic Reasoning** have been proposed. This framework integrates external, specialized agents as tools within the main reasoning process. It identifies three universally effective agentic tools: a **Web-Search agent** for information retrieval, a **Coding agent** for quantitative computations, and a **Mind-Map agent** to serve as a structured memory.<sup>40</sup> The Mind-Map agent is a key innovation for maintaining coherence in long reasoning chains. It dynamically constructs a knowledge graph from the reasoning context, extracting entities and their relationships. This structured representation can then be queried by the main reasoning LLM, allowing it to retrieve critical information and maintain logical consistency over long and complex tasks.<sup>40</sup>

## Diagnosing Failures in Agentic Workflows

When a complex, multi-step agentic workflow fails, a systematic approach is needed to diagnose the root cause efficiently. A recommended best practice is a two-phase evaluation process<sup>42</sup>:

1. **End-to-End Task Success Evaluation:** The first phase treats the agent as a black box. The primary goal is to determine if the agent successfully achieved the user's ultimate goal. This is typically measured with a binary success/failure metric, often assessed by a human or a highly capable LLM judge. This phase quickly identifies which workflows are failing.
2. **Step-Level Diagnostics:** Once a failure is identified, the second phase involves "opening the box" to analyze the agent's internal trajectory. Assuming the system is well-instrumented with detailed logs of all actions, tool calls, and intermediate reasoning, each step of the failed workflow can be scored independently. Key checkpoints for evaluation include:
  - **Tool Choice:** Was the correct tool selected for the given sub-task?
  - **Parameter Extraction:** Were the inputs for the tool call complete, well-formed, and correctly extracted from the context?
  - **Error Handling:** If a tool returned an error or an empty result, did the agent recover gracefully and attempt a valid alternative action?<sup>33</sup>

To identify systemic patterns of failure across many interactions, a **Transition Failure Matrix** is an exceptionally powerful diagnostic tool. This matrix visualizes where failures most commonly occur in a workflow. The rows of the matrix represent the last successful state or step, and the columns represent the state where the first failure occurred. By aggregating data from many failed runs, this matrix immediately highlights the "hotspots" in the workflow—the transitions that are most fragile and responsible for the highest number of failures. This data-driven approach allows development teams to move beyond anecdotal debugging of individual traces and focus their engineering efforts on the parts of the system that will yield the greatest improvements in overall reliability.<sup>42</sup>

## Strategic Recommendations for Enterprise-Grade Agent Deployment

Deploying LLM and multi-agent systems into production requires a strategic,

engineering-driven approach to evaluation that extends far beyond academic leaderboards. A successful strategy must be continuous, holistic, and deeply integrated into the development lifecycle. The following recommendations provide a framework for building and maintaining reliable, high-quality agentic systems in an enterprise environment.

## A Checklist for Production-Ready Continuous Automated Testing

To ensure consistent quality and prevent regressions, teams should implement a continuous automated testing pipeline based on the following checklist:

- **Establish a "Golden" Evaluation Dataset:** Curate a version-controlled dataset of high-quality test cases that are representative of both common user interactions and critical edge cases. This dataset serves as the ground truth for regression testing.<sup>35</sup>
- **Define a Portfolio of Metrics:** Select and implement a suite of evaluation metrics that provide a holistic view of agent performance. This should include:
  - **Behavioral Metrics:** Task success rate, output correctness.
  - **Capability Metrics:** Tool call accuracy, planning coherence.
  - **Quality Metrics:** Factual consistency, relevance, absence of toxicity.
  - **Operational Metrics:** Latency, token cost, resource usage.<sup>20</sup>
- **Integrate Evaluation into CI/CD Pipelines:** Automate the execution of your evaluation suite within your CI/CD platform (e.g., GitHub Actions, GitLab CI). Configure pipelines to run automatically on triggers such as pull requests to the main branch or on a nightly schedule.<sup>19</sup>
- **Set Automated Regression Thresholds:** Define clear, quantitative thresholds for your key metrics. Configure the CI/CD pipeline to automatically fail the build if a change causes performance to drop below these thresholds, creating an effective quality gate.<sup>19</sup>
- **Implement Comprehensive Observability and Tracing:** Instrument the agentic system to log every step of its execution, including reasoning traces, tool calls, parameters, and responses. This detailed tracing is a prerequisite for effective debugging and advanced, offline evaluation.<sup>33</sup>
- **Establish a Human-in-the-Loop (HITL) Review Process:** For failures that are complex, nuanced, or occur in high-stakes scenarios, automated evaluation may be insufficient. Establish a clear workflow for escalating these failures to human experts for review. The insights from these reviews should be used to improve the evaluation dataset and automated scorers.<sup>33</sup>

## Selecting and Combining Benchmarks for a Holistic Evaluation

## Strategy

No single benchmark can capture all dimensions of an agent's performance. A robust evaluation strategy must therefore consist of a carefully selected portfolio of benchmarks tailored to the specific requirements of the system being built. The following decision framework can guide the selection process:

- **If the agent's primary function is code generation:**
  - Start with **LiveCodeBench** to establish a baseline for generalization and to test for contamination against older benchmarks.
  - Incorporate **AutoCodeBench** to assess performance across a broad and diverse set of programming languages, especially if multilingual support is a requirement.
  - Add **SWE-Bench** if the agent is expected to perform complex, real-world software engineering tasks at the repository level.
- **If the system involves multiple agents coordinating on complex plans:**
  - **REALM-Bench** is essential. It is specifically designed to test the core multi-agent capabilities of planning, coordination, and adaptation to dynamic disruptions.
- **If the agent will be deployed to a global, multilingual user base:**
  - **BenchMAX** is the critical benchmark. It provides the most comprehensive framework for evaluating advanced capabilities like reasoning, tool use, and long-context understanding across a wide range of languages, ensuring equitable performance for all users.
- **For a baseline of general reasoning and problem-solving ability:**
  - **AGIEval** offers a stable, human-grounded measure of a model's general cognitive skills, providing a useful signal of its foundational intelligence.

By combining these specialized benchmarks, teams can construct a comprehensive evaluation suite that provides a multi-faceted and accurate picture of their agent's strengths and weaknesses.

## Future Directions: The Next Frontier of Evaluation

The field of LLM and agent evaluation is evolving at a breakneck pace. As agentic systems become more sophisticated and autonomous, evaluation methodologies must continue to advance to keep pace. The next frontier of evaluation will likely focus on several key areas:

- **More Holistic and Realistic Benchmarks:** There is a clear need for evaluation environments that more closely simulate the complexity and unpredictability of the real world. This includes developing benchmarks that test an agent's ability to learn and adapt over long periods, manage complex memory structures, and interact seamlessly in

multimodal environments.<sup>3</sup>

- **Standardized Evaluation of Multi-Agent Systems:** While REALM-Bench provides a strong foundation, the evaluation of multi-agent collaboration is still a nascent field. The proliferation of agentic frameworks like CrewAI, LangGraph, and AutoGen will necessitate the development of standardized benchmarks and metrics for assessing inter-agent communication, coordination strategies, and emergent collaborative behaviors.<sup>45</sup>
- **Scalable and Automated Safety and Alignment Verification:** As agents gain more autonomy, ensuring they remain safe and aligned with human values becomes paramount. Future research will need to focus on developing scalable, automated methods for verifying agent safety, identifying potential harmful behaviors, and ensuring compliance in complex, open-ended scenarios.

Ultimately, the continuous evolution of evaluation frameworks is the critical engine driving progress in AI. By building better tools to measure, verify, and diagnose our systems, we enable the development of more capable, reliable, and trustworthy AI agents that can be deployed with confidence to solve real-world problems.

## Cytowane prace

1. Evaluation and Benchmarking of LLM Agents: A Survey - arXiv, otwierano: września 22, 2025, <https://arxiv.org/html/2507.21504v1>
2. Evaluation and Benchmarking of LLM Agents: A Survey - arXiv, otwierano: września 22, 2025, <https://arxiv.org/pdf/2507.21504>
3. [2507.21504] Evaluation and Benchmarking of LLM Agents: A Survey - arXiv, otwierano: września 22, 2025, <https://arxiv.org/abs/2507.21504>
4. LiveCodeBench: Holistic and Contamination Free Evaluation of Large Language Models for Code, otwierano: września 22, 2025, <https://livecodebench.github.io/>
5. AutoCodeBench: Large Language Models are Automatic ... - arXiv, otwierano: września 22, 2025, <https://arxiv.org/abs/2508.09101>
6. AutoCodeBench: Large Language Models are Automatic Code Benchmark Generators, otwierano: września 22, 2025, <https://arxiv.org/html/2508.09101v1>
7. AutoCodeBench: Large Language Models are Automatic Code Benchmark Generators, otwierano: września 22, 2025, <https://autocodebench.github.io/>
8. Tencent-Hunyuan/AutoCodeBenchmark - GitHub, otwierano: września 22, 2025, <https://github.com/Tencent-Hunyuan/AutoCodeBenchmark>
9. ruixiangcui/AGIEval - GitHub, otwierano: września 22, 2025, <https://github.com/ruixiangcui/AGIEval>
10. REALM-Bench: A Benchmark for Evaluating Multi-Agent Systems on Real-world, Dynamic Planning and Scheduling Tasks - arXiv, otwierano: września 22, 2025, <https://arxiv.org/html/2502.18836v2>
11. [2502.18836] REALM-Bench: A Benchmark for Evaluating Multi-Agent Systems on Real-world, Dynamic Planning and Scheduling Tasks - arXiv, otwierano: września 22, 2025, <https://arxiv.org/abs/2502.18836>
12. REALM-Bench: A Benchmark for Evaluating Multi-Agent ... - arXiv, otwierano:



- września 22, 2025, <https://arxiv.org/pdf/2502.18836>
13. BenchMAX: A Comprehensive Multilingual Evaluation Suite for Large Language Models, otwierano: wrzesnia 22, 2025, <https://arxiv.org/html/2502.07346v1>
  14. CONE-MT/BenchMAX - GitHub, otwierano: wrzesnia 22, 2025, <https://github.com/CONE-MT/BenchMAX>
  15. BenchMAX: A Comprehensive Multilingual Evaluation Suite for Large Language Models, otwierano: wrzesnia 22, 2025, <https://arxiv.org/html/2502.07346v2>
  16. [Literature Review] BenchMAX: A Comprehensive Multilingual Evaluation Suite for Large Language Models - Moonlight, otwierano: wrzesnia 22, 2025, <https://www.themoonlight.io/en/review/benchmax-a-comprehensive-multilingual-evaluation-suite-for-large-language-models>
  17. Transform CI/CD Pipeline: Harness Automated Code Insights - Medium, otwierano: wrzesnia 22, 2025, <https://medium.com/@API4AI/transform-ci-cd-pipeline-harness-automated-code-insights-92e35733f200>
  18. Ultimate guide to CI/CD: Fundamentals to advanced implementation - GitLab, otwierano: wrzesnia 22, 2025, <https://about.gitlab.com/blog/ultimate-guide-to-ci-cd-fundamentals-to-advanced-implementation/>
  19. Continuous Evaluation of Generative AI Using CI/CD Pipelines - WillowTree, otwierano: wrzesnia 22, 2025, <https://www.willowtreeapps.com/craft/continuous-evaluation-of-generative-ai-using-ci-cd-pipelines>
  20. Real-Time LLM Evaluation: Continuous Testing for Production AI ..., otwierano: wrzesnia 22, 2025, [https://medium.com/@future\\_agi/real-time-llm-evaluation-continuous-testing-for-production-ai-7bc7df3ee437](https://medium.com/@future_agi/real-time-llm-evaluation-continuous-testing-for-production-ai-7bc7df3ee437)
  21. Actions · LiveCodeBench/LiveCodeBench - GitHub, otwierano: wrzesnia 22, 2025, <https://github.com/LiveCodeBench/LiveCodeBench/actions>
  22. Automate your project with GitHub Models in Actions, otwierano: wrzesnia 22, 2025, <https://github.blog/ai-and-ml/generative-ai/automate-your-project-with-github-models-in-actions/>
  23. Continuous Benchmark · Actions · GitHub Marketplace, otwierano: wrzesnia 22, 2025, <https://github.com/marketplace/actions/continuous-benchmark>
  24. Implement a GitHub Actions workflow for code formatting and testing - Tilburg Science Hub, otwierano: wrzesnia 22, 2025, <https://tilburgsciencehub.com/topics/automation/automation-tools/deployment/github-actions-workflow/>
  25. GitHub Actions, otwierano: wrzesnia 22, 2025, <https://github.com/features/actions>
  26. CI/CD YAML syntax reference - GitLab Docs, otwierano: wrzesnia 22, 2025, <https://docs.gitlab.com/ci/yaml/>
  27. Writing .gitlab-ci.yml File with Examples [Tutorial] - Spacelift, otwierano: wrzesnia 22, 2025, <https://spacelift.io/blog/gitlab-ci-yml>
  28. GitLab CI/CD examples, otwierano: wrzesnia 22, 2025,



- <https://docs.gitlab.com/ci/examples/>
29. What Is the .gitlab-ci.yml File and How to Work With It - Codefresh, otwierano: września 22, 2025,  
<https://codefresh.io/learn/gitlab-ci/what-is-the-gitlab-ci-yml-file-and-how-to-work-with-it/>
  30. A collection of gitlab-ci examples - GitHub, otwierano: września 22, 2025,  
<https://github.com/dokku/gitlab-ci>
  31. Evaluate Langfuse LLM Traces with an External Evaluation Pipeline ..., otwierano: września 22, 2025,  
[https://langfuse.com/guides/cookbook/example\\_external\\_evaluation\\_pipelines](https://langfuse.com/guides/cookbook/example_external_evaluation_pipelines)
  32. How to evaluate an LLM evaluation framework - Vellum AI, otwierano: września 22, 2025,  
<https://www.vellum.ai/blog/how-to-evaluate-an-llm-evaluation-framework>
  33. How to Evaluate AI Agents: Comprehensive Strategies for Reliable, High-Quality Agentic Systems - Maxim AI, otwierano: września 22, 2025,  
<https://www.getmaxim.ai/articles/how-to-evaluate-ai-agents-comprehensive-strategies-for-reliable-high-quality-agentic-systems/>
  34. Evaluating AI Agents in 2025: A Practical Guide - Turing College, otwierano: września 22, 2025,  
<https://www.turingcollege.com/blog/evaluating-ai-agents-practical-guide>
  35. LLM Evaluation: Frameworks, Metrics, and Best Practices | SuperAnnotate, otwierano: września 22, 2025,  
<https://www.superannotate.com/blog/llm-evaluation-guide>
  36. !! Top 5 Open-Source LLM Evaluation Frameworks in 2025 - DEV Community, otwierano: września 22, 2025,  
<https://dev.to/guybuildingai/-top-5-open-source-llm-evaluation-frameworks-in-2024-98m>
  37. Verifying Chain-of-Thought: Labeling Reasoning Steps in Model Outputs - Labelvisor, otwierano: września 22, 2025,  
<https://www.labelvisor.com/verifying-chain-of-thought-labeling-reasoning-steps-in-model-outputs/>
  38. Multi-Step Reasoning with Large Language Models, a Survey - arXiv, otwierano: września 22, 2025, <https://arxiv.org/html/2407.11511v2>
  39. lz1oceani/verify\_cot - GitHub, otwierano: września 22, 2025,  
[https://github.com/lz1oceani/verify\\_cot](https://github.com/lz1oceani/verify_cot)
  40. A Streamlined Framework for Enhancing LLM ... - ACL Anthology, otwierano: września 22, 2025, <https://aclanthology.org/2025.acl-long.1383.pdf>
  41. A Streamlined Framework for Enhancing LLM Reasoning with Agentic Tools - ACL Anthology, otwierano: września 22, 2025,  
<https://aclanthology.org/2025.acl-long.1383/>
  42. Q: How do I evaluate agentic workflows? - Hamel's Blog, otwierano: września 22, 2025,  
<https://hamel.dev/blog/posts/evals-faq/how-do-i-evaluate-agentic-workflows.html>
  43. How to Evaluate LLM Applications: The Complete Guide - Confident AI,

otwierano: września 22, 2025,

<https://www.confident-ai.com/blog/how-to-evaluate-llm-applications>

44. Evaluating Agentic AI Workflows - The Couchbase Blog, otwierano: września 22, 2025, <https://www.couchbase.com/blog/evaluating-agentic-ai-workflows/>
45. The Leading Multi-Agent Platform, otwierano: września 22, 2025, <https://www.crewai.com/>
46. Top 5 Open-Source Agentic Frameworks - Research AIMultiple, otwierano: września 22, 2025, <https://research.aimultiple.com/agentic-frameworks/>
47. Comparative Analysis of LLM Agent Frameworks | by Jose F. Sosa - Medium, otwierano: września 22, 2025, <https://medium.com/@josefsosa/white-paper-comparative-analysis-of-llm-agent-frameworks-3d9ea8c0212f>
48. A Comprehensive Guide to Evaluating Multi-Agent LLM Systems - Orq.ai, otwierano: września 22, 2025, <https://orq.ai/blog/multi-agent-llm-eval-system>